



Ghulam Ishaq Khan Institute of Engineering Sciences and Technology
Faculty of Computer Science & Engineering

MicroLend

Community Microfinance Loan Servicing and Delinquency Engine

SE423 — Software Construction and Development

Milestone 1 — Proposal, Object-Oriented Design and Pattern Plan

Submitted by

Muhammad Ibrahim	2023446
Hassan Khalid	2023242
Tughrul Hussain	2023532

Instructor: Dr. Mian Muaz Razaq

Fall 2026

Contents

MicroLend — Community Microfinance Loan Servicing & Delinquency Engine	1
1. Problem Definition and Requirement-to-Minimum Mapping	1
2. Module Breakdown and Ownership	3
3. Object-Oriented Design	4
4. SOLID Application Plan	9
5. Design Pattern Selection and Justification	11
6. Repository and Documentation Setup	13
7. Declaration	14
Annex A — AI Usage Log	15
Milestone 1	15
Milestone 2	17
Milestone 3	17
Rules the team has agreed for AI use on this project	17

MicroLend — Community Microfinance Loan Servicing & Delinquency Engine

Team

Member	Reg. #	Primary modules	Pattern owned
Muhammad Ibrahim	2023446	origination, delinquency	Adapter
Hassan Khalid	2023242	products, reporting	Factory Method
Tughral Hussain	2023532	ledger	Decorator

Team Lead / Integrator (M1): Tughral Hussain — role rotates each milestone. **Repository:** <https://github.com/tughrall1/microlend>
Milestone: M1 (Week 6) · **Weight:** 15% of the project grade

1. Problem Definition and Requirement-to-Minimum Mapping

1.1 The system

A community microfinance institution lends small amounts — typically Rs. 25,000 to Rs. 200,000 — to shopkeepers, tradespeople and smallholders who cannot access bank credit. The loans are short, repayment is frequent, and the institution's survival depends on knowing, every single day, exactly how much is owed, how much is overdue, and how much of the portfolio is at risk of never being recovered.

MicroLend is the back-office engine that answers those questions. It assesses loan applications against eligibility rules, generates repayment schedules, applies incoming repayments to the right obligations in the right order, maintains an audit-grade ledger, and runs a nightly process that ages overdue accounts, accrues penalties and classifies the portfolio by risk.

The system is a **back-office engine with a small web interface that runs locally**. It is never deployed: there is no hosting, no external service and no paid resource of any kind. The application serves its pages from a local process against a file-based database, so every demonstration runs entirely on a team member's own laptop.

1.2 Its users

User	What they do with MicroLend
Loan officer	Enters an application; the system approves or rejects it with a stated reason, and produces the repayment schedule the customer signs
Teller / cashier	Posts repayments as customers pay; the system decides how each payment is applied and records it
Branch manager	Reads portfolio reports — how much is overdue, how much must be provisioned, how effective collection has been
System (scheduled)	Runs the end-of-day process: accrues penalties, ages accounts, re-classifies risk, writes off unrecoverable loans

1.3 Why this problem is worth building software for

Three properties make this domain demanding in a way a CRUD application is not:

1. **Money must be exactly right.** A repayment schedule whose instalments do not sum to principal plus interest is not “approximately correct”; it is wrong, and the error compounds across a portfolio.
2. **Order of application matters.** The same Rs. 4,000 applied penalty-first versus principal-first produces different balances, different interest, and a different delinquency classification months later.
3. **The record must be trustworthy.** A ledger that can be edited is not a ledger. Corrections must be visible as corrections.

Those three properties are what give the system its invariants, and the invariants are what the design in Section 3 protects.

1.4 Scope boundary

In scope for this project: individual loans, three product types, fixed-rate lending in a single currency, repayment by cash at the branch or by bank statement import, penalty accrual, delinquency classification, provisioning, write-off, statements and portfolio reporting.

Explicitly deferred (recorded here so that scope cannot creep silently): group/joint-liability lending, guarantors and collateral registers, loan restructuring and rescheduling, multi-currency, savings products, interest on savings, mobile money integration, and any deployment of the application beyond running it locally.

1.5 Mapping of each Section 3.1 minimum requirement to how it will be met

#	Minimum requirement (guidelines §3.1)	How MicroLend meets it	Evidence at
1	≥ 4 distinct functional modules with real business logic, not a single-file CRUD demo	Five modules: origination (eligibility decisions), products (schedule generation), ledger (allocation waterfall and posting), delinquency (penalty accrual, ageing, provisioning), reporting (statements and portfolio metrics). Each owns a distinct decision , not a screen — see §2	M2, M3
2	OO design demonstrating SOLID , with written justification of where and why each principle was applied	§3 gives the class design with cohesion and coupling reasoning; §4 analyses all five principles, identifying where each is relevant, not relevant, or not yet observable	M1 (§3, §4)
3	≥ 3 correctly implemented, justified patterns , at least one creational plus one structural or behavioural	Factory Method (creational) in products, Adapter (structural) at the bank-statement import boundary, Decorator (structural) on the repayment-posting pipeline. Each with the alternative rejected — §5	M1 plan, M2 implementation
4	≥ 1 genuine concurrent or asynchronous operation with a documented thread-safety strategy	The end-of-day engine processes the active portfolio across a fixed thread pool while tellers post repayments to the same accounts. Strategy in §3.6	M2

#	Minimum requirement (guidelines §3.1)	How MicroLend meets it	Evidence at
5	Defensive programming and structured exception handling across module boundaries	Validation in constructors and command factories so invalid objects cannot exist; a typed exception hierarchy rooted at <code>LoanDomainException</code> ; <code>Money</code> rejects invalid values at construction — §3.5	M2
6	Persistent data layer with a defined schema	SQLite, seven tables with foreign keys and CHECK constraints that re-state the code's invariants at the storage layer — §3.7	M2
7	TDD suite $\geq 60\%$ of core business logic with visible red-green-refactor history	Eligibility, schedule generation, allocation and classification are pure functions over <code>Money</code> and dates — ideal test-first targets. Commit convention <code>test(red):</code> → <code>feat(green):</code> → <code>refactor::</code> squash-merge disabled in repository settings so the trail survives — §6.3	Repository history, M2, M3
8	Hosted Git repository with individually attributable contribution history	GitHub, public. Module ownership maps one-to-one onto directories, so <code>git log</code> per path is meaningful. CODEOWNERS routes each pull request to the module owner; no author merges their own work — §6	Repository, all milestones
9	CI pipeline that builds and runs the test suite on push	GitHub Actions: compile, JUnit, JaCoCo coverage report, Checkstyle gate. No secrets, no network dependency — §6.4	M2

2. Module Breakdown and Ownership

Five modules — one above the minimum, so that if effort runs short, reporting can be reduced without dropping below four.

#	Module	Scope in one line	Primary owner
1	origination	Decides whether an application is approved, on what terms, and with what stated reason if rejected	Muhammad Ibrahim
2	products	Holds the loan product catalogue and generates the repayment schedule for an approved loan	Hassan Khalid
3	ledger	Applies incoming repayments through the allocation waterfall and records every movement in an append-only ledger	Tughrul Hussain
4	delinquency	Runs the end-of-day process: accrues penalties, computes days past due, classifies and provisions accounts, writes off	Muhammad Ibrahim

#	Module	Scope in one line	Primary owner
5	reporting	Produces customer statements, portfolio risk reports and CSV exports	Hassan Khalid

2.1 The business rules each module owns

These are frozen at M1 in the **Business Rules Register** (docs/business-rules.md) so that the requirements the SQA side traces against do not move underneath it.

origination — An applicant must be 18 or older at disbursement and no older than 65 at maturity; may hold at most 2 active loans; total monthly instalment burden must not exceed 40% of declared monthly income; must have no prior write-off; the requested amount and tenure must lie within the product's bands. A failing application yields a **typed rejection reason**, not a boolean. Risk grade A/B/C adjusts the rate by a fixed spread and caps the approved amount.

products — Three products: **flat-rate** (interest on the original principal throughout), **reducing-balance** (interest on the outstanding balance, equal instalments), and **bullet** (interest monthly, principal at maturity). Instalment dates follow the same day-of-month as disbursement, clamped for short months. An optional grace period defers the first instalment but still accrues interest. All rounding residue is absorbed into the final instalment, so the instalments sum **exactly** to principal plus total interest.

ledger — Payments are applied in strict order: **penalty** → **fees** → **interest** → **principal**, oldest unpaid instalment first. Partial payments are allowed and applied as far as they reach. An overpayment becomes an advance credit, unless it clears the loan, in which case early settlement applies a rebate of unearned interest for reducing-balance loans and no rebate for flat-rate loans. **A posted ledger entry is immutable**; a mistake is corrected only by a compensating reversal entry.

delinquency — For a given business date: accrue a penalty of 0.05% per day on overdue principal, capped at 10% of the overdue amount; compute days past due; place the account in a bucket (0 / 1–30 / 31–60 / 61–90 / 90+) mapping to a classification (Standard / Watch / Substandard / Doubtful / Loss); apply the matching provision percentage; write the loan off automatically past 180 days overdue.

reporting — Customer statement for a date range; portfolio reports (portfolio at risk over 30 days, provision coverage, collection efficiency); CSV export; an alert event raised whenever an account's classification changes.

2.2 Why ownership is split this way

Each owner holds a module whose **invariant they can state in a single sentence**, which is what makes individual Q&A defensible:

- **Muhammad Ibrahim** — *“No loan may exist that the eligibility rules would not have approved.”*
- **Hassan Khalid** — *“The instalments of a schedule sum exactly to principal plus total interest.”*
- **Tughrul Hussain** — *“The ledger always balances, and nothing in it is ever altered.”*

Ibrahim additionally owns the concurrent end-of-day engine, which depends on Tughrul's ledger; that dependency is deliberate, because it forces the two of them to agree an interface rather than reach into each other's internals.

SE431 deliverable ownership rotates so that every member leads each kind of quality artifact at least once:

Milestone	SQAP & risk register	Quality requirements & RTM	Inspection, static analysis & white-box
M1	Muhammad Ibrahim	Hassan Khalid	Tughrul Hussain
M2	Hassan Khalid	Tughrul Hussain	Muhammad Ibrahim
M3	Tughrul Hussain	Muhammad Ibrahim	Hassan Khalid

3. Object-Oriented Design

3.1 Design intent


The design is organised around one idea: **the objects that hold money-related state are responsible for defending their own correctness**. Validation is not performed by services on passive data structures; it happens inside the objects, at construction and at every state transition, so that an invalid object cannot exist for other code to defend against.

Everything else follows from that. Services are thin and stateless. Data crosses module boundaries as immutable value objects. Persistence is reached only through interfaces the domain itself declares.

3.2 Class diagram

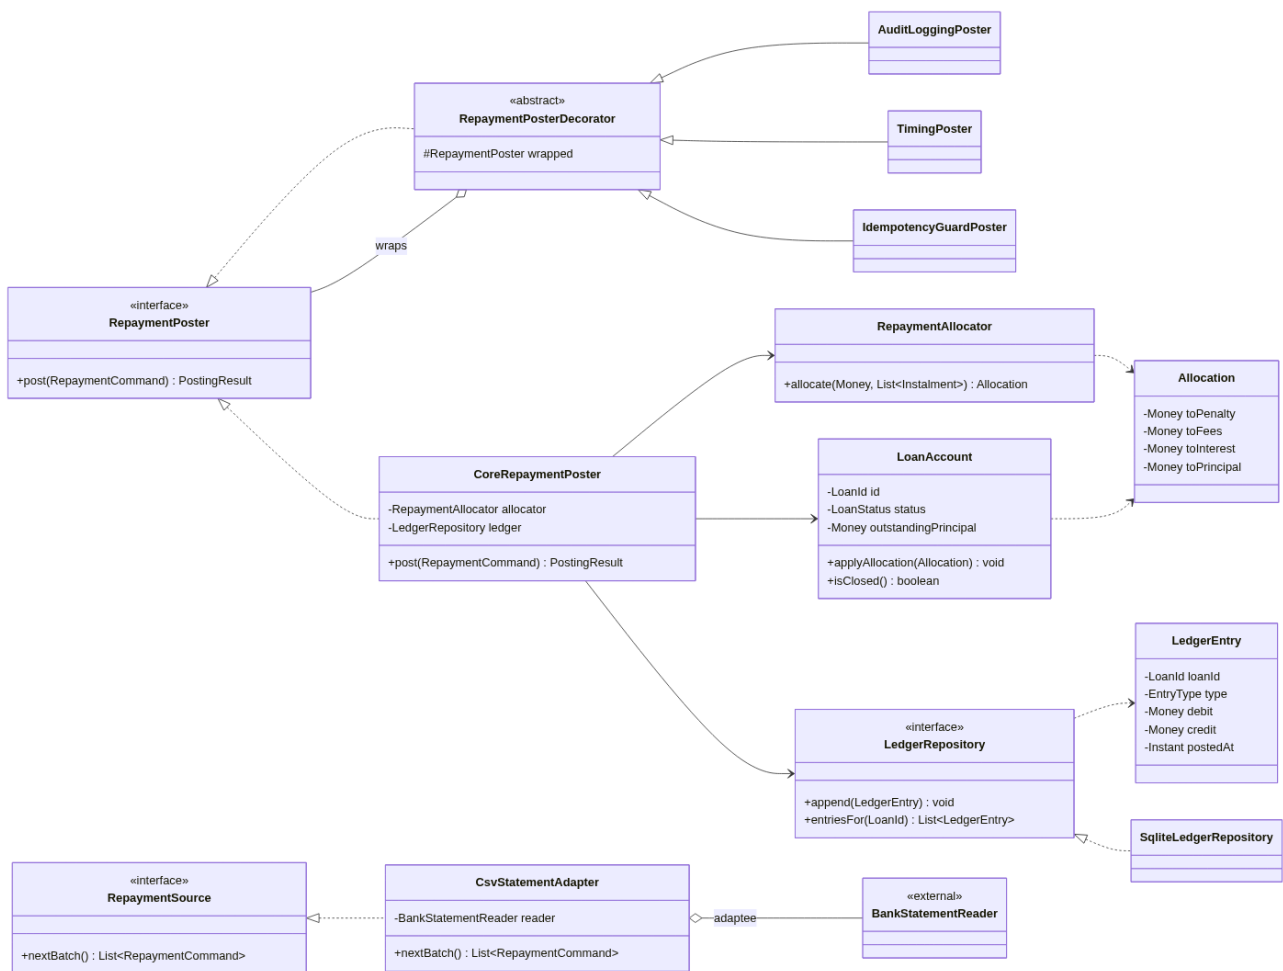
The class model is shown in three views, one per area of the system. Shared kernel types (Money, LoanId, the exception hierarchy) appear where they are used.

View 1 — Origination and Products: deciding whether to lend, and on what schedule



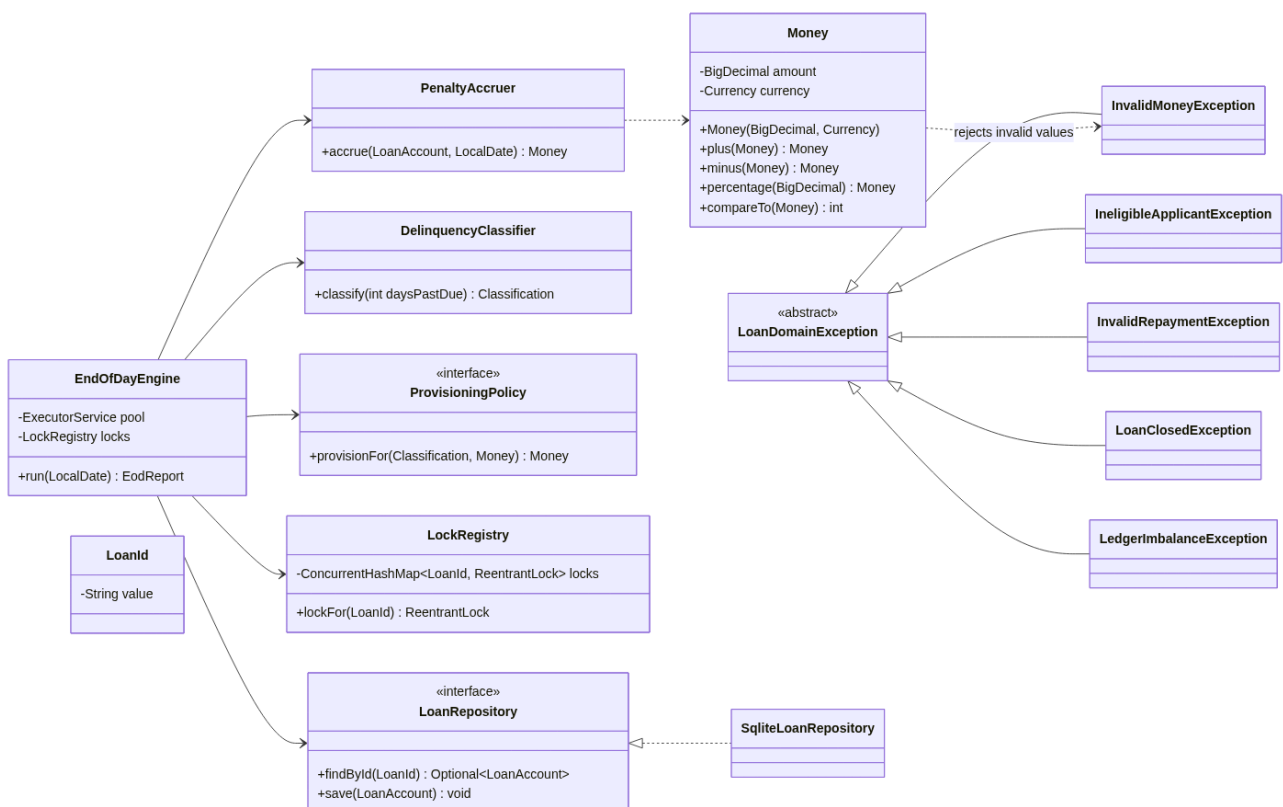
Note the two extension points: a new eligibility rule is a new *EligibilityRule*; a new loan product is a new *LoanProduct* subclass with its calculator. Neither requires editing a tested class.

View 2 — Ledger: applying money, with the posting pipeline and the import boundary




The Decorator chain and the Adapter boundary are both visible here: decorators are each IS-A and HAS-A Repayment-Poster, and CsvStatementAdapter holds the foreign reader by composition rather than inheriting it.

View 3 — Delinquency, the shared kernel, and the persistence boundary



EndOfDayEngine holds only the interfaces it needs — it cannot export a report or read the full ledger, because those methods are not on the types it is given. The exception hierarchy is rooted at *LoanDomainException*, so nothing outside a module ever catches a storage-layer exception.

3.3 Component view and allowed dependencies



The dependency rule: arrows point inward, toward common, and never outward from a domain module to persistence or the web layer. persistence depends on the domain, not the reverse — the dashed arrows are implementations of interfaces the domain declares. There are **no cycles** between modules, and this is enforced as a build-time check rather than left as an intention.

3.4 Cohesion and coupling reasoning

Cohesion — each module is functionally cohesive, not grouped by convenience.

The test applied to every module was: *do these operations exist to serve one job, or do they merely share a noun?* A class holding `assessEligibility()`, `saveToDatabase()`, `renderStatementPdf()` and `emailCustomer()` would be only **logically cohesive** — grouped because each mentions a loan — and would answer to four different stakeholders. That class does not exist here. Instead:

Module	Its single job	What was deliberately kept out
origination	Decide approval and terms	Schedule arithmetic (that is products), persistence (that is an injected repository)
products	Turn an approved amount and tenure into a schedule	Deciding <i>whether</i> to lend, and anything about repayment
ledger	Apply money to obligations and record the movement	Deciding penalties or classification (that is delinquency), formatting output
delinquency	Age, penalise, classify and provision	Writing ledger lines itself — it asks ledger to do that
reporting	Present what has already happened	Any calculation that changes state

Within modules the same test is applied to classes. *EligibilityEvaluator* holds no rules of its own; each rule is a separate *EligibilityRule* implementation, so adding “must not be a staff member’s relative” adds a class rather than an

if. `RepaymentAllocator` computes an `Allocation` and does not write it; `CoreRepaymentPoster` writes it and does not compute it.

Coupling — loose, and in one direction.

1. **No module constructs another module's dependencies.** Business classes never execute new `SQLiteLoanRepository(...)`. Repositories arrive through constructors, so `ledger` depends on the `LedgerRepository` *interface it declares itself*, and the `SQLite` implementation depends inward on that interface. The effect is concrete: the posting logic can be tested with an in-memory fake, with no database on the machine, and swapping `SQLite` for `PostgreSQL` changes one class in persistence and nothing in any domain module.
2. **Data crosses boundaries as immutable value objects.** `Money`, `Instalment`, `LedgerEntry` and `Allocation` have no setters. A caller cannot mutate another module's state by holding a reference to something it was given, and no module needs defensive copying on receipt.
3. **No internal collection is ever returned.** `Schedule.instalments()` returns an unmodifiable view. Returning the live list would let a caller bypass every rule in `Schedule`, and would freeze the choice of `List` as a public API that could never be changed to a `Map` for lookup by due date.
4. **Content coupling is impossible by construction.** No module exposes mutable fields; all state is private and changed only through methods that enforce the invariant.
5. **The one accepted dependency with weight** is `delinquency` → `ledger`, because penalty accrual must write `ledger` entries. It is deliberately **interface-only**: the end-of-day engine holds a `LedgerWriter`, not the posting implementation, and cannot read statements or export CSV through it.

Where coupling is highest, and why that is acceptable: `ledger` is the most depended-upon domain module (`delinquency` and reporting both need it). This is *afforent* coupling — many things depend on it — which is appropriate for a stable module whose contract is the smallest and least likely to change in the system. The risk this creates is recorded as **R1** in the risk register, and its mitigation is the golden-schedule test suite written before the implementation.

3.5 Defensive programming and exception design

Invalid objects cannot be constructed.

```
public final class Money implements Comparable<Money> {
    private final BigDecimal amount;

    public Money(BigDecimal amount, Currency currency) {
        if (amount == null)          throw new IllegalArgumentException("amount is required");
        if (amount.scale() > 2)      throw new InvalidMoneyException("scale exceeds 2: " + amount);
        if (amount.signum() < 0)     throw new InvalidMoneyException("negative amount: " + amount);
        this.amount = amount.setScale(2, RoundingMode.HALF_EVEN);
        this.currency = Objects.requireNonNull(currency, "currency is required");
    }
}
```

Because `Money` cannot hold a negative or over-scaled value, no downstream code tests for one. The same applies to `LoanApplication` (validated in its constructor) and `RepaymentCommand` (validated in its factory). This is the difference between defending against bad input at one boundary and scattering `if (amount < 0)` through five modules.

A typed exception hierarchy, carrying the reason.

```
LoanDomainException (abstract – the module boundary type)
├─ IneligibleApplicantException (carries the failing RuleOutcome)
├─ InvalidRepaymentException   (carries the offending value and its source)
├─ LoanClosedException         (carries loan id and closure date)
├─ LedgerImbalanceException    (carries the computed debit and credit totals)
└─ InvalidMoneyException
```

Rules: no catch (`Exception e`), and no empty catch blocks — both are rejected by the static-analysis configuration. Exceptions crossing a module boundary are always `LoanDomainException` subtypes, so a caller never catches a `JDBC` exception; the repository translates storage failures into domain terms. Each exception carries the data needed to explain the failure, because `throw new RuntimeException("invalid")` tells an operator nothing.

3.6 The concurrent operation and its thread-safety strategy

What runs concurrently. The end-of-day engine processes the active portfolio across `Executors.newFixedThreadPool(4)`, one task per loan account, while teller threads are still posting repayments to the same accounts. This is not contrived:

in a real branch, the nightly batch and late counter transactions overlap, and the demo harness runs both at once.

The three races it must prevent:

1. **Lost update.** A teller reads an outstanding balance of 10,000; the end-of-day task reads 10,000; both write back. One movement vanishes and the ledger stops balancing.
2. **Double penalty accrual.** If the engine runs twice for the same business date — a retry after failure, or two pool threads taking the same account — penalty is charged twice, inventing money.
3. **Torn read.** A portfolio report sums balances mid-update and reports a total that was never true at any instant.

The strategy:

Measure	What it prevents
Immutability at the core — Money, LedgerEntry, Instalment have no mutable state	Nothing shared can be changed under another thread
Per-account striped locks — <code>ConcurrentHashMap<LoanId, ReentrantLock></code> ; all balance-changing work happens under that account's lock, never held across I/O	Lost updates, while leaving unrelated accounts fully parallel
Ascending LoanId lock ordering for any operation touching two accounts	Deadlock, by a stated and inspectable rule
Unique (loan_id, business_date) in <code>delinquency_snapshot</code> as an idempotency key	Double accrual — the second attempt is rejected by the schema even if locking fails
Append-only ledger , with an invariant test asserting $\Sigma \text{debits} = \Sigma \text{credits}$ after every run	Silent imbalance under any interleaving
Deterministic tests driven by <code>CountDownLatch</code> , never <code>Thread.sleep</code>	Flaky CI, which is what causes teams to start ignoring red builds

This also supplies the two deliberately injected failure scenarios the M3 rubric asks for: a repository that throws mid-posting must leave the ledger balanced, and an end-of-day task that dies must not block the pool or leave a half-applied accrual.

3.7 Persistent data layer

SQLite, seven tables. The schema **re-states the code's invariants** so that a corrupt write fails at two independent levels:

Table	Holds	Constraints that matter
customer	Applicant identity, income, write-off history	—
loan_product	Catalogue: code, rate, amount band, tenure band	CHECK (min_amount <= max_amount)
loan	Account state: principal, status, grade, dates	FK to customer and product; CHECK (outstanding >= 0)
instalment	Generated schedule lines	FK to loan; UNIQUE (loan_id, sequence)
ledger_entry	Every money movement	FK to loan; CHECK (debit >= 0 AND credit >= 0); no UPDATE statement exists anywhere in the codebase
payment	Received payments and their source	UNIQUE (bank_reference) — the idempotency guard
delinquency_snapshot	End-of-day result per loan per date	UNIQUE (loan_id, business_date) — prevents double accrual

Monetary columns are stored as INTEGER minor units (paisa) rather than REAL, for the same reason double is banned in the code.

4. SOLID Application Plan

The rubric requires all five principles to be analysed, each marked **relevant**, **not relevant**, or **not yet observable**, with non-applicability justified. Three are demonstrated through concrete design decisions below.

4.1 Single Responsibility Principle — relevant, demonstrated

The test applied: who can demand a change to this class? If two different stakeholders can, it has two reasons to change.

The design decision this produced: an early sketch had one `LoanService` holding eligibility checks, schedule generation, repayment posting and statement rendering. Four actors could demand changes to it — the credit committee (eligibility), the finance department (rates and schedules), the branch operation (how payments are applied) and the marketing/communications side (statement layout). It was split into `EligibilityEvaluator`, `LoanProduct/ScheduleGenerator`, `RepaymentAllocator` and `StatementRenderer`.

The concrete consequence: a change to the penalty rate touches delinquency alone; a change to statement layout touches reporting alone. Neither forces re-testing of the repayment waterfall. Had they stayed together, a statement-formatting change would require re-certifying money-handling code.

`LoanAccount` holds state and enforces its own invariants, and does nothing else — it neither computes schedules nor renders itself.

4.2 Open–Closed Principle — relevant, demonstrated

The design decision: adding a fourth loan product must not require editing any tested class.

`LoanProduct` defines the stable `generateSchedule()` workflow and declares an abstract `createInterestCalculator()`. A new product is a new subclass plus a new calculator — two new files. `generateSchedule()` is never reopened, so the three existing products' tests cannot be broken by the arrival of a fourth.

The same shape appears in origination: eligibility is a list of `EligibilityRule` objects, so a new rule is a new class and a registration line, not an edit to `EligibilityEvaluator`.

What was rejected: a switch (`productType`) inside schedule generation. It would place every product's arithmetic in one method that must be reopened — and re-tested in full — for every new product, with the risk that a mistake in the new branch changes what existing borrowers are charged.

4.3 Liskov Substitution Principle — relevant, demonstrated

The design decision, stated as something deliberately *not* done: the institution offers interest-free hardship loans. The tempting modelling is `HardshipLoan` extends `ReducingBalanceProduct` with `createInterestCalculator()` overridden to throw `UnsupportedOperationException`, since there is no interest to calculate.

That is rejected. It refuses an operation the supertype promises, so a `HardshipLoan` cannot stand in for a `LoanProduct`, and every caller would need `instanceof` to avoid the exception. Instead, hardship lending is a `ZeroInterestCalculator` that honours the contract and returns `Money.ZERO`. Substitution holds, and no caller changes.

The contract every `InterestCalculator` honours: same preconditions (any non-negative principal, any tenure within the product band); returns a non-negative `Money`, never null; a schedule always has exactly `tenureMonths` instalments. These are stated in the interface's Javadoc and asserted in a **shared contract test** that every implementation must pass — so LSP is checked by the build rather than by good intentions.

4.4 Interface Segregation Principle — relevant, demonstrated

The design decision: the persistence boundary is not one fat `LoanDao`.

Separate role interfaces: `ScheduleProvider` (read a schedule), `LedgerWriter` (append an entry), `LedgerReader` (query entries), `RepaymentPoster` (post a repayment), `ReportExporter` (write CSV).

The concrete consequence: the end-of-day engine receives `ScheduleProvider` and `LedgerWriter` only. It is therefore *incapable* of exporting a CSV or reading the full ledger — not by policy, but because those methods are not on the interfaces it holds. A fat interface would have made the engine's test double implement export methods it never calls, and every one of those stubs would be a latent failure waiting to be invoked.

4.5 Dependency Inversion Principle — relevant, demonstrated

The design decision: the domain declares the interfaces it needs; persistence implements them.

`LoanRepository` and `LedgerRepository` are declared **inside the domain modules**, not in the persistence package. `SqliteLoanRepository` implements them and is injected at the composition root — the one place in the application that is allowed to call new on infrastructure.

The concrete consequence: no business class imports `java.sql`. Unit tests inject in-memory fakes and run with no database present. A ledger-posting test runs in milliseconds and cannot fail because of a locked file or a missing driver.

What was rejected: new `SqliteLoanRepository(connection)` inside `RepaymentPoster`. That would make the posting logic untestable without a real database, push a JDBC concern into business code, and make a storage change a business-code change.

4.6 Summary and honest limits

Principle	Status at M1	Where it is demonstrated
SRP	Relevant, demonstrated	Splitting of <code>LoanService</code> into evaluator, generator, allocator, renderer — §4.1
OCP	Relevant, demonstrated	<code>LoanProduct</code> hierarchy and the <code>EligibilityRule</code> list — §4.2
LSP	Relevant, demonstrated	Hardship lending modelled as a zero calculator, plus the shared contract test — §4.3
ISP	Relevant, demonstrated	Role interfaces at the persistence and posting boundary — §4.4
DIP	Relevant, demonstrated	Domain-declared repository interfaces, injected implementations — §4.5

Not yet observable at M1. All five are present in the design, but three can only be *evidenced* once code exists: LSP’s contract test must run against three real calculators (M2); ISP’s benefit shows when a test double is written for each role interface (M2); DIP’s payoff is visible when the in-memory fakes make the test suite run without a database (M2). M1 commits to the structure; M2 produces the evidence.

Where a principle is weaker, stated honestly. SRP at the *module* level is clean, but `RepaymentAllocator.post()` is a single method with roughly eleven branches, which is the most complex method in the system. It is kept as one method deliberately — splitting the waterfall across classes would scatter a rule that must be read in order — and it is the primary subject of the Fagan inspection and the basis-path test design on the SQA side. Complexity that is acknowledged, inspected and fully covered is a different proposition from complexity that is hidden.

5. Design Pattern Selection and Justification

Three patterns, chosen by the rule that the construction problem comes first and the pattern name second. One creational and two structural, satisfying the “at least one creational plus at least one structural or behavioural” requirement.

5.1 Factory Method — `LoanProduct` as Creator (creational)

Owner: Hassan Khalid

The construction problem. `generateSchedule()` is a stable workflow: validate the tenure, build the due-date series, compute each instalment’s interest, absorb the rounding residue into the final instalment, assemble the schedule. Exactly **one** step varies — which interest calculation applies. If the workflow writes new `ReducingBalanceCalculator()`, the schedule module becomes coupled to concrete classes through the `new` keyword, and must be reopened for every new product.

The structure. `LoanProduct` is the **Creator** and holds the stable workflow. `createInterestCalculator()` is the **factory method**, declared abstract by the Creator and overridden by `FlatRateProduct`, `ReducingBalanceProduct` and `BulletProduct` — the **ConcreteCreators**. `InterestCalculator` is the **Product** abstraction, and the three calculators are **ConcreteProducts**. The creator subtype determines what is created; the workflow never names a concrete class.

```
public abstract class LoanProduct {
    public final Schedule generateSchedule(Money principal, int tenureMonths, LocalDate start) {
        InterestCalculator calculator = createInterestCalculator(); // ← the factory method
        // ... stable workflow: dates, per-instalment interest, rounding residue, assembly
    }
    protected abstract InterestCalculator createInterestCalculator();
}
```

Alternative rejected: Simple Factory. A `ScheduleFactory.create(type)` holding a switch was the obvious first design. It is rejected because the creation decision stays centralised in one place that must be edited for every new product — it *moves* the conditional ladder rather than removing it. A method that returns a new object is not automatically the GoF Factory Method; the distinguishing property is that the **creator subtype**, not a parameter, decides what gets created.

Also rejected: Abstract Factory, because there is no product *family* here — only one axis of variation, so a second factory layer is indirection without a second dimension. **Also rejected: Singleton** for a shared product registry, because it hides dependencies, introduces global mutable state and leaks state between tests; constructor injection already gives one shared instance without the coupling.

5.2 Adapter — CsvStatementAdapter (structural)

Owner: Muhammad Ibrahim

The construction problem. Repayments also arrive as bank statement files, and every bank's format differs: column order, date format, and amount convention — several report amounts in **minor units** (paisa) as integers. Neither contract is wrong; they were designed independently. The servicing core must keep thinking in `Money` and `LocalDate`.

The structure. The **Client** is the repayment-posting workflow. The **Target** is `RepaymentSource`, the interface our application already understands. The **Adaptee** is each bank's statement reader. The **Adapter** — `HblStatementAdapter`, `McbStatementAdapter` — converts 254900 into `Money(2549.00)` and 04102026 into a `LocalDate`, and translates the reader's parse failures into `InvalidRepaymentException`. It is an **Object Adapter**: it *has-a* reader rather than inheriting one.

Alternative rejected: Facade. The distinction is intent: an Adapter protects clients from an **interface mismatch**; a Facade protects clients from **excessive subsystem knowledge**. At the import boundary there is no complicated subsystem to simplify — there is exactly one foreign contract that does not fit ours. That is mismatch, not complexity.

Also rejected: changing the core to speak the bank's format. That moves the dependency in the wrong direction, spreading vendor vocabulary — minor units, their date format, their sign convention — through business logic, which is precisely the coupling the Adapter exists to contain at one boundary.

5.3 Decorator — the RepaymentPoster pipeline (structural)

Owner: Tughrul Hussain

The construction problem. Posting a repayment must sometimes also (a) write an audit journal line naming the operator, (b) record elapsed time for the GQM measurement programme, and (c) guard against posting the same bank reference twice. **Which of these are needed differs per call path, and they combine:**

Call path	Audit	Timing	Idempotency
Teller screen	✓	✓	—
CSV bulk import	one per file	✓	✓
End-of-day internal adjustment	—	✓	—
Unit tests	—	—	—

The structure. `RepaymentPoster` is the **Component**; `CoreRepaymentPoster` is the concrete component holding the real allocation call; `RepaymentPosterDecorator` both **IS-A** `RepaymentPoster` and **HAS-A** `RepaymentPoster`. `AuditLoggingPoster`, `TimingPoster` and `IdempotencyGuardPoster` each add one focused behaviour and then delegate to `wrapped.post(c)`.

Composition order is deliberate and documented: the idempotency guard wraps outermost, so a duplicate is rejected before anything is journalled or timed. Order matters in a Decorator chain, and stating why is part of the design.

```
RepaymentPoster tellerPoster =
    new AuditLoggingPoster(
        new TimingPoster(
            new CoreRepaymentPoster(allocator, ledgerRepository)));
```

Alternative rejected: inheritance combinations. `AuditingPoster`, `TimingPoster`, `AuditingTimingPoster`, `AuditingIdempotentPoster`... — eight classes for three optional behaviours, and sixteen the day a retry behaviour is added. This is the classic subclass explosion, and the combinations cannot be chosen at runtime per call path.

Also rejected: Proxy. The two have nearly the same structure and are separated by intent: a Decorator primarily **extends behaviour**; a Proxy primarily **controls access**. Here nothing is being gated, protected or stood in for remotely — optional responsibilities are being layered, and they must stack. That is Decorator's pressure, not Proxy's.

5.4 Patterns considered and deliberately not used at M1

Pattern	Why not now
Facade — <code>RepaymentFacade.postRepayment()</code>	A genuine candidate: three clients need the same five-step workflow. Held as a possible fourth pattern at M2. If adopted, it must coordinate and not absorb — the allocation rules stay in <code>RepaymentAllocator</code> , or it becomes a god class
Strategy — <code>ProvisioningPolicy</code>	Valid, and likely at M2 for swapping provisioning percentages when regulations change. Not claimed at M1, and <code>InterestCalculator</code> will not be presented as a Strategy, because it is the Product of the Factory Method above; claiming both patterns for the same objects would be incoherent
State — loan lifecycle; Observer — classification alerts	Both fit the domain, and both are deferred to M2 rather than adopted on vocabulary the team cannot yet defend in depth

A pattern is only included where it solves a problem that existed in the design *before* the pattern was named. The three above each do.

6. Repository and Documentation Setup

Repository: <https://github.com/tughrall/microlend> (public) **Collaborators:** tughrall, mIBRAHIM707, Hassan242-kk

6.1 Branching strategy

Three tiers:

Branch	Purpose	Rules
main	Milestone baselines only	Protected. No direct pushes. Changes arrive solely by reviewed pull request
develop	Integration branch, the current working state	Receives pull requests from feature branches
feature/<module>-<short-name>	One unit of work, owned by one member	Branched from develop, merged back by pull request

6.2 Quality gates (configured, not aspirational)

A GitHub **ruleset** is active on main:

- Direct pushes are **blocked** — every change requires a pull request
- **One approving review** required, from someone other than the author
- Stale approvals are **dismissed** when new commits are pushed
- **Last-push approval** required, so code cannot be slipped in after approval
- All review conversations must be **resolved** before merge
- Branch deletion and force-pushes are blocked
- Only the repository owner may bypass, and every bypass is logged by GitHub

CODEOWNERS routes each pull request automatically to the owner of the module it touches. Shared code — Money, the exception hierarchy, the Checkstyle configuration — requires all three members.

6.3 Preserving the TDD trail

Squash-merge and rebase-merge are disabled in repository settings; merge commits are the only permitted method.

This is deliberate: squashing collapses a red-green-refactor sequence into a single commit, and the evidence cannot be reconstructed afterwards.

Commit convention:

```
test(red):    instalment rounding residue lands on the final instalment
feat(green):  absorb rounding residue in ScheduleGenerator
refactor:     extract RoundingPolicy from ScheduleGenerator
```

The pull-request template requires the author to name the red, green and refactor commit hashes for the change.

6.4 Baselines and CI

Milestone submissions are tagged on main: `m1-baseline`, `m2-baseline`, `m3-baseline`. A tag is created only at submission, so it points at exactly the state that was assessed.

The CI pipeline (GitHub Actions) runs on every push: `compile` → `JUnit` → `JaCoCo coverage report` → `Checkstyle gate`. It requires no secrets and no network access.

6.5 Coding standards

`config/checkstyle.xml` is the project coding standard, and it enforces **constraint C-01: money is never a binary floating-point type**. `double` and `float` are rejected by the build, and a second rule catches stray decimal literals such as `0.05`. All monetary values use `BigDecimal` inside `Money` with an explicit `RoundingMode`. `0.1 + 0.2 != 0.3` in IEEE 754, and a fraction of a rupee per instalment accumulates across a portfolio into a ledger that does not balance.

The configuration also enforces complexity limits (cyclomatic complexity ≤ 15 per method, classes ≤ 400 lines), rejects empty catch blocks and `catch (Exception)`, and forbids public mutable fields.

6.6 Documentation

File	Contents
<code>README.md</code>	System overview, module table with owners, build and run instructions, branching strategy, commit convention, coding standards
<code>docs/M1-SE423-*.md</code>	This document
<code>docs/business-rules.md</code>	The Business Rules Register — every rule frozen at M1, with a deferred list
<code>qa/risk-register.md</code>	Living risk register, initiated at M1 with 8 scored risks
<code>qa/ai-usage-log.md</code>	AI Usage Log, submitted with every milestone
<code>config/checkstyle.xml</code>	The coding standard, referenced by SQAP §4.4

7. Declaration

The design in this document is the team’s own. Where AI assistance was used — in drafting documentation and in reviewing design alternatives — it is recorded in `qa/ai-usage-log.md` with the tool, purpose, what was adopted, what was modified, and how it was verified. Each member is the owner of their modules and is able to explain and defend the design decisions recorded here.

Annex A — AI Usage Log

Required by the joint-project guidelines §9 (Academic Integrity & AI Use), submitted with every milestone.

The guidelines permit AI-assisted development, but assessment is based on whether the team understands, verifies and can defend what it submits. Any member who cannot explain or justify work claimed as theirs may receive reduced individual marks regardless of the team's score.

How this log is maintained: one row per significant use, added **at the time of use**, not reconstructed before a deadline. The *Verification* column names what was actually checked.

Milestone 1

Entry 1 — Select a project that satisfies both course rubrics

Date: 4 October 2026 · **Tool:** Claude Opus 5 (Claude Code)

Activity. Asked for candidate projects meeting every Section 3.1 and 3.2 minimum, simple to build, and defensible in a viva. Three candidates were analysed and scored against the rubric.

Adopted. The comparative analysis in `project/Project-Idea-Analysis.md`; MicroLend selected over SkyHold (air-line seat inventory) and DispenseRx (pharmacy dispensing).

Modifications made. Scope reduced to an application that runs locally with no deployment, no hosted services and no paid resources. DispenseRx rejected by the team on the stated ground that its clinical data would be invented and therefore indefensible under questioning.

Verification performed. Fit table checked row by row against the guidelines PDF, with each claimed minimum traced to where it is met. Team reviewed and accepted the recommendation on 4 October 2026.

Entry 2 — Draft the SE423 M1 document

Date: 4 October 2026 · **Tool:** Claude Opus 5 (Claude Code)

Activity. Asked for a document following the six M1 rubric criteria for the chosen project.

Adopted. Draft of `docs/M1-SE423-Proposal-00D-Pattern-Plan.md`: class model, SOLID analysis, pattern justifications and the cohesion/coupling reasoning.

Modifications made. The single wide class diagram was replaced by three per-area views after the first PDF proved illegible. “Command-line interface” was replaced throughout with a locally served web interface, once the team confirmed the intended form of the application.

Verification performed. Pattern choices checked against the taught material. An earlier draft's “Factory Method” was found to be a Simple Factory — which Lecture 11 states is *not* the GoF pattern — and was restructured so that `LoanProduct` subclasses are the `ConcreteCreators`. Team reviewed and accepted the document on 4 October 2026.

Entry 3 — Draft the SE431 M1 document

Date: 4 October 2026 · **Tool:** Claude Opus 5 (Claude Code)

Activity. Asked for quality requirements, an RTM, the SQAP outline, Cost of Quality, standards awareness and a risk register, against the SE431 M1 rubric.

Adopted. Draft of `docs/M1-SE431-Quality-Requirements-RTM-SQAP.md`.

Modifications made. Usability and portability entries rewritten for a locally-run web application. Quality requirements restated so each names one of the four standard verification methods, and four items reclassified from quality requirements to constraints (C-01 to C-04).

Verification performed. Requirement targets and the Cost of Quality estimate reviewed for realism against the team's available time. Team reviewed and accepted the document on 4 October 2026.

Entry 4 — Repository, quality gates and coding standard

Date: 4 October 2026 · **Tool:** Claude Opus 5 (Claude Code)

Activity. Asked to create the repository with branch protection, CODEOWNERS, a pull-request template, and a Checkstyle configuration enforcing the money constraint.

Adopted. `config/checkstyle.xml`, `.github/CODEOWNERS`, `.github/pull_request_template.md`, and the README branching section.

Modifications made. The C-01 check was rewritten from a text pattern to a token-based rule (`IllegalToken` on `NUM_DOUBLE` and `NUM_FLOAT`) after CI showed the original flagging the version string `"0.1.0-M1"` as a decimal literal.

Verification performed. Branch protection tested by attempting a direct push to `main`, which GitHub reported as a bypassed rule violation. The Checkstyle fault was found by the CI pipeline itself and corrected before the build went green.

Entry 5 — Business Rules Register

Date: 4 October 2026 · **Tool:** Claude Opus 5 (Claude Code)

Activity. Asked to turn the described rules into a numbered register with boundary values for test design.

Adopted. `docs/business-rules.md` — 43 rules across the five modules.

Modifications made. None to the rule set. The six numeric decisions listed in §8 — the 40% burden ceiling, the 0.05%/day penalty capped at 10%, the 180-day write-off, the risk-grade spreads, the provision percentages and the 3-month grace cap — were reviewed and **accepted unchanged** by the team on 4 October 2026.

Verification performed. Boundary values cross-checked against the test design techniques recorded in the RTM, so every rule carrying a threshold has a corresponding boundary test planned.

Entry 6 — Build, test scaffold and CI pipeline

Date: 4 October 2026 · **Tool:** Claude Opus 5 (Claude Code)

Activity. Asked for a Gradle build, the first domain class written test-first, and a CI pipeline, with nothing that incurs cost.

Adopted. `build.gradle`, the Gradle wrapper, `Money`, `LoanDomainException`, `InvalidMoneyException`, `MoneyTest`, and `.github/workflows/build.yml`.

Modifications made. Written as a genuine TDD cycle: the failing test was committed first (`test(red):`), then the implementation (`feat(green):`). `Money.minus` was given a negative-result guard so the type cannot express direction — the ledger's debit and credit columns do that.

Verification performed. CI run on the public repository: compile, 10 JUnit tests, the Checkstyle gate and the coverage report all pass. Two faults were found and fixed through CI — a missing executable bit on `gradlew`, and the Checkstyle false positive described in entry 4.

Entry 7 — Submission PDFs

Date: 4 October 2026 · **Tool:** Claude Opus 5 (Claude Code)

Activity. Asked for a PDF build with a GIKI title page carrying the team names, project title and course code.

Adopted. `docs/build-pdf.sh` and `docs/header.tex`.

Modifications made. Fonts changed to Times New Roman with explicit Unicode mappings.

Verification performed. Rendered pages inspected individually. A font fault was found that silently dropped the symbols for “less than or equal to”, “greater than or equal to”, “approximately” and “multiplied by” from the PDFs — which would have corrupted the stated quality targets — and was corrected before submission.

Declaration for Milestone 1

I confirm that I have read, understood and verified the material submitted under my name, that I can explain and defend the design decisions, quality requirements and analysis it contains, and that any AI-assisted content was reviewed and modified by me where needed.

Muhammad Ibrahim — 2023446 Sections owned: origination, delinquency; SQAP outline, risk register

Signature: Date:

Hassan Khalid — 2023242 Sections owned: products, reporting; quality requirements, RTM

Signature: Date:

Tughral Hussain — 2023532 Sections owned: ledger; standards awareness, Cost of Quality, repository and build setup

Signature: Date:

Milestone 2

(rows added as work proceeds)

#	Date	Tool	Purpose	Prompt / activity summary	Adopted	Modifications	Verification
---	------	------	---------	---------------------------------	---------	---------------	--------------

Milestone 3

(rows added as work proceeds)

#	Date	Tool	Purpose	Prompt / activity summary	Adopted	Modifications	Verification
---	------	------	---------	---------------------------------	---------	---------------	--------------

Rules the team has agreed for AI use on this project

1. **Generated code enters the repository only through the module owner**, who reviews it line by line and rewrites what they would have written differently. The owner is the one answering for it in the viva.
2. **No generated test is accepted without reading what it asserts.** A test that passes for the wrong reason is worse than no test.
3. **Business rules and numeric thresholds are the team’s decisions**, not suggestions accepted by default.
4. **Every row in this log is added when the work happens.** A log assembled the night before a deadline is itself evidence of the problem the requirement is designed to catch.
5. **If a member cannot explain a submitted artifact, it is withdrawn or rewritten before submission** rather than defended in the demo.